


Cloud Resource Allocation — Greedy BFD (baseline) vs Simulated Annealing [performance focus]
Simulated Annealing improves F(X) by 24.2% · 36 of 50 tasks moved



Greedy BFD (baseline)

F(X) = 1.2608 energy 11,921 W latency 38,313 ms servers 5/10 feasible



Simulated Annealing

F(X) = 0.9552 energy 13,156 W latency 26,681 ms servers 10/10 feasible



Low priority Medium priority High priority Memory gauge (fill = used / capacity) Idle server (0 W) Task on a different server than under Greedy BFD (baseline)